

Nonlinear Machine Learning and Ensemble Methods

Prof. Mingkui Tan


SCUT Machine Intelligence Laboratory (SMIL)



SMIL内部资料 请勿外泄

Classification Examples

- How to classify?



SMIL内部资料 请勿外泄

Contents

1 Decision Tree

2 Introduction of Ensemble Learning

3 AdaBoost

4 GBDT (Gradient Boosting Decision Trees)

Contents

1 Decision Tree

2 Introduction of Ensemble Learning


3 AdaBoost

4 GBDT(Gradient Boosting Decision Trees)

SMIL内部资料 请勿外泄

Classification Examples

- How to classify?



SMIL内部资料 请勿外泄

Decision Tree Example

■ Learning the ‘Play Tennis Decision Tree’

4 Attributes

Day	Outlook	Temp	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

Decision Tree Example

Play Tennis: Yes or No?

- Learned function is a tree
- Each branch of the tree represents a possible decision, outcome, or reaction
- The farthest branches (leaf nodes) on the tree represent the end results


Algorithm: Decision Tree

Input: Training set $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$; **Attribute set** $A = \{a_1, a_2, \dots, a_d\}$.

Procedure: Function TreeGenerate(D, A). **Output:** A decision tree.

```
1 Generate node;  
2 if all the samples in  $D$  belong to class  $C$  then  
3   mark this node as class  $C$  leaf node; return  
4 end  
5 if  $A = \emptyset$  OR samples in  $D$  have the same value on  $A$  then  
6   mark this node as leaf node, and the class should be the most frequent  
   occurrence class; return  
7 end  
8 Select the best partition attribute  $a_*$  from  $A$ ;  
9 for each value  $a_*^v$  of attribute  $a_*$  do  
10   $D_v$  is the sample subset of  $D$  with  $a_* = a_*^v$ ;  
11  if  $D_v = \emptyset$  then  
12    mark this node as the leaf node, and the class should be the most frequent  
    occurrence class; return  
13  else  
14    generate a branch for this node, TreeGenerate( $D_v, A \setminus \{a_*\}$ )  
15  end  
16 end
```

Decision Tree Example



Which is the best partition attribute?

SMIL内部资料 请勿外泄

Decision Tree Example

An attribute is good when:

- For one value we get all instances as positive
- For other value we get all instances as negative

An attribute is poor when:

- It provides no discrimination
- Attribute is immaterial to the decision
- Same number of positive and negative instances for each value

Measure of Homogeneity of Examples

- **Entropy** characterizes the (im)purity of an arbitrary collection of examples
- Given a collection D of positive and negative examples
- **Entropy of D** relative to boolean classification is

$$\text{Entropy}(D) = - p_+ \log_2 p_+ - p_- \log_2 p_-$$

where p_+ is proportion of positive examples

p_- is proportion of negative examples

Measure Homogeneity of Examples

Illustration:


- D is a collection of 14 examples with 9 positive and 5 negative examples
- Entropy of D relative to the boolean classification:

$$\text{Entropy}(9+, 5-) = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.940$$

- Entropy is zero if all members of D belong to the same class

Measure Homogeneity of Examples

Entropy Function Relative to a Boolean Classification



Entropy for Multi-valued Target function

- If the target attribute can take on c different values, the entropy of D relative to this **c-wise** classification is:

$$\text{Entropy}(D) = \sum_{i=1}^c -p_i \log_2 p_i$$

Information Gain


- Entropy measures the (im)purity of a collection
- Information gain of attribute A is the reduction in entropy caused by partitioning the set of examples D

$$\text{Gain}(D, A) \equiv \text{Entropy}(D) - \sum_{v \in \text{Values}(A)} \frac{|D_v|}{|D|} \text{Entropy}(D_v)$$

where $\text{Values}(A)$ is the set of all possible values for attribute A

D_v is the subset of D for which attribute A has value v

Measure of Purity: Information



Info[2,3]=0.97bits

Info[4,0]=0bits


Info[3,2]=0.97bits

$$\text{Info}[2,3] = \text{entropy}(2,3) = -\frac{2}{5} \log_2 \frac{2}{5} - \frac{3}{5} \log_2 \frac{3}{5} = 0.97 \text{ bits}$$

Information Gain for Each Attribute

- $\text{Gain}(\text{outlook}) = 0.94 - 0.693 = 0.247$
- $\text{Gain}(\text{temperature}) = 0.94 - 0.911 = 0.029$
- $\text{Gain}(\text{humidity}) = 0.94 - 0.788 = 0.152$
- $\text{Gain}(\text{windy}) = 0.94 - 0.892 = 0.048$
- $\text{argmax}_A\{0.247, 0.029, 0.152, 0.048\} =$
outlook
- Select outlook as the splitting attribute of tree


Expanded Tree: Play Tennis for Outlook = Sunny



- Where $temp$ contains wind, humidity and temperature


SMIL内部资料 请勿外泄

Expanded Tree: Play Tennis for Outlook = Sunny




SMIL内部资料 请勿外泄

Expanded Tree: Play Tennis for Outlook = Sunny



- Since $\text{Gain}(\text{humidity})$ is the highest, select humidity as splitting attribute

Decision Tree for Weather Data



SMIL内部资料 请勿外泄

Contents

1 Decision Tree

2 Introduction of Ensemble Learning


3 AdaBoost

4 GBDT (Gradient Boosting Decision Trees)

SMIL内部资料 请勿外泄


Ensemble Learning

- Ensemble learning: combine several base models to produce a better predictive model
- Main methods: Bagging, Boosting



Bagging

- Get T sampling sets through bootstrap sampling
- Train T base learners through the sampling sets respectively



- For classification:
The class with the most votes becomes the final class
- For regression:
The final output is the average output of every base learner

Random Forest


- Random forest is an extension of bagging using decision tree as base learner
- Randomly select m out of p features to get the optimal partition feature

Comparison between bagging and random forest

- The training efficiency of random forest is better than bagging
- Bagging uses decision tree with **definite structure**
- Random forest uses decision tree with **random structure**

Random Forest

- An example of the process flow is depicted below:



SMIL内部资料 请勿外泄

Contents

1 Introduction of Ensemble Learning

2 Decision Tree

3 AdaBoost

4 GBDT (Gradient Boosting Decision Trees)


SMIL内部资料 请勿外泄

AdaBoost: What is AdaBoost?

- Stands for Adaptive Boosting
- Combines base learners linearly
- Iteratively adapts to the errors made by base learners in previous iterations
- Re-weighting scheme:
 - Higher weight assigned to incorrectly classified data points
 - Lower weight assigned to correctly classified data points

AdaBoost: How does it work?

- Learns a series of base classifiers from the data
- Combines base learners linearly
- Makes strong predictions based on combined learners



Algorithm: AdaBoost

Input: Training set $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$; Weak learning algorithm.

Output: A strong classifier.


- 1 Initialize the weight vector: $w_1(i) = 1/n, i = 1, 2, \dots, n$;
- 2 for $t = 1, \dots, T$ do
- 3 fit base learner $h_t(x) \in \{-1, +1\}$
- 4 Calculate the classification error rate e_t of $h_t(x)$
$$e_t = p(h_t(x) \neq y_i) = \sum_{i=1}^N w_t(i) \mathbb{I}(h_t(x_i) \neq y_i)$$
where $\mathbb{I}\{\cdot\}$ is the indicator function: $\mathbb{I}(X = x_i) = \begin{cases} 1, & h_t(x_i) \neq y_i \\ 0, & h_t(x_i) = y_i \end{cases}$
- 5 Calculate the weight α_t of $h_t(x)$
- 6 Update the weights of each data point

$$\alpha_t = \frac{1}{2} \ln \frac{1 - e_t}{e_t}$$

- 7
- 8 end
- 9 Output the final hypothesis: $H(\mathbf{x}) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(\mathbf{x}))$

AdaBoost: Initialization of the algorithm

Training data



- Assigning same weight $w_1(i)$ to all data points

$$w_1(i) = \frac{1}{n}, i = 1, 2, \dots, n$$


where n is the number of data points

AdaBoost: Iterations of Algorithm

For $t = 1, \dots, T$ or until low enough error is achieved:

- Fit base learner $h_t(x)$ to data points
- Calculate the classification error rate e_t of $h_t(x)$
- Calculate the weight α_t of $h_t(x)$
- Update the weights of each data point


AdaBoost: Iteration 1




- Initially all data points have the same weight $1/n$
- The class correctly classified will be given less weights in the next iteration, and higher weights for misclassified class

AdaBoost: Iteration 2


Weights after 1st iteration



Fit base learner
 $h_2(x)$




Calculate weights




- Base learner forms a decision boundary which classifies the data points better with higher weights

AdaBoost: Iteration 3


Weights after 2nd iteration



Fit base learner
 $h_3(x)$



Calculate weights




...continue iterating until either:

- Sufficiently low training error is achieved (with enough iterations, the algorithm can reach 100% accuracy)
- A pre-defined number of base learners was added

AdaBoost: Final Model


$$H(x) = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t(x)\right)$$

base learner $h_1(x)$




α_1

base learner $h_2(x)$




$+ \alpha_2$

base learner $h_3(x)$



$+ \alpha_3$

strong model $H(x)$



=

AdaBoost: More Details

- Fit base learner $h_t(x)$ to data points, we define base classifier $h_t(x)$ as

$$h_t(x) \in \{-1, +1\}$$

- Calculate the classification error rate e_t of $h_t(x)$ on the training data set

$$e_t = p(h_t(x) \neq y_i) = \sum_{i=1}^n w_t(i) \mathbb{I}(h_t(x_i) \neq y_i)$$

For binary classification case, $\mathbb{I}(X = x_i) = \begin{cases} 1, & h_t(x_i) \neq y_i \\ 0, & h_t(x_i) = y_i \end{cases}$

AdaBoost: More Details

- Calculate the weight α_t of $h_t(x)$ in the final classifier

$$\alpha_t = \frac{1}{2} \ln \frac{1 - e_t}{e_t}$$

Here, when $e_t \leq 0.5$, $\alpha_t \geq 0$, with the decreasing of e_t , α_t will be larger

Note: classifier with smaller error rate would be more important

AdaBoost: More Details

- Update the weights

$$w_{t+1}(i) = \frac{w_t(i)}{z_t} e^{-\alpha_t y_i h_t(x_i)}$$

Here, $z_t = \sum_{i=1}^n w_t(i) e^{-\alpha_t y_i h_t(x_i)}$ is normalization term,

and makes $w_t(i)$ a probability distribution

To simplify, $w_{t+1}(i) = \begin{cases} \frac{w_t(i)}{z_t} e^{-\alpha_t}, & \text{for right predictive sample} \\ \frac{w_t(i)}{z_t} e^{\alpha_t}, & \text{for wrong predictive sample} \end{cases}$

AdaBoost: More Details


- Ensemble learner

$$H(\mathbf{x}) = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t(\mathbf{x})\right)$$

Note: $h_t(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$ is a nonlinear function, so AdaBoost can deal with nonlinear problem

AdaBoost: More Details

Every iteration generates a new base learner $h_t(x)$ with weight α_t



AdaBoost: Example

Given the following training samples:

Index	1	2	3	4	5	6	7	8	9	10
x	0	1	2	3	4	5	6	7	8	9
y	-1	-1	-1	1	1	1	-1	-1	-1	1

- Initialize the same weights for data points at the beginning, $w_1(1), w_1(2), \dots, w_1(10) = 0.1$
- Assume the threshold of x is 2.5, when $x \leq 2.5, h_1(x) = -1$, when $x > 2.5, h_1(x) = 1$
- So the right samples are Index 1, 2, 3, 4, 5, 6, 10, and the wrong samples are Index 7, 8, 9 (three data points)
- The error rate $e_1 = 3 * 0.1 = 0.3$
- For all samples, the threshold is chosen to be 2.5, which has the lowest error rate e_1

AdaBoost: Example

- Now, the first base classifier is

$$h_1(x) = \begin{cases} 1, & x \leq 2.5 \\ -1, & x > 2.5 \end{cases}$$

- Then, for $h_1(x)$, α_1 is

$$\alpha_1 = \frac{1}{2} \ln \frac{1 - e_1}{e_1} = 0.424$$

- For the first round, the weight of $h_1(x)$ is 0.424, it is used to update the weights of data points of the next round

$$w_{t+1}(i) = \frac{w_t(i)}{z_t} e^{-\alpha_t y_i h_t(x_i)}$$

- $w_2(i)$, $i = 1, \dots, 10$ are as shown in the table below:

index	1	2	3	4	5	6	7	8	9	10
x	0	1	2	3	4	5	6	7	8	9
y	-1	-1	-1	1	1	1	-1	-1	-1	1
w_2	0.0714	0.0714	0.0714	0.0714	0.0714	0.0714	0.1667	0.1667	0.1667	0.0714

AdaBoost: Why Use AdaBoost?

- Needs only a simple classifier as a base learner
- Can achieve prediction similar to powerful classifiers
- Can combine with any learning algorithms
- Requires little parameter tuning
- Extended to problems beyond binary classification

AdaBoost: Summary

- AdaBoost (i.e. Adaptive boosting) is one of the most popular and powerful ensemble methods
- AdaBoost focuses on the data points that are erroneous
- Simple to implement and depend on the base learner
- However, vulnerable to noisy data

Contents

1 Introduction of Ensemble Learning

2 Decision Tree

3 AdaBoost

4 GBDT (Gradient Boosting Decision Trees)

SMIL内部资料 请勿外泄

Gradient Boosting Decision Trees

Gradient Boosting Decision Trees (GBDT) is a decision tree algorithm with iteration

Example: What is the difference between regression tree and GBDT?

Suppose: There are 4 people A, B, C and D, whose age are 14, 16, 24, 26 respectively





Figure: Single regression tree

GBDT

- The key of GBDT is that trees learn all the results and residuals of all trees before
- The residual is the difference of predictive value and real value, so the predictive value is the sum of all results of trees

For GBDT



- So, $A = 15 + (-1) = 14$, $B = 15 + 1 = 16$, $C = 25 + (-1) = 24$, $D = 25 + 1 = 26$

Gradient Boosting Decision Trees: Questions

Q1: Why do we need GBDT?

- The motivation of this algorithm
 - Every calculation of residual is to increase the weight of wrong predictive samples
 - The residual of right predictive sample is zero
- In the next iteration, the model can concentratively address these wrong predictive samples
- Another function is to prevent overfitting

Gradient Boosting Decision Trees: Questions

Q2: Where does this algorithm reflect gradient boosting?

- Residual is the gradient descent direction, which is the derivation of mean square error (MSE)
- MSE is the loss function of CART regression tree

Setting	Loss Function	$-\partial L(y_i, f(x_i))/\partial f(x_i)$
Regression	$\frac{1}{2}[y_i - f(x_i)]^2$	$y_i - f(x_i)$
Regression	$ y_i - f(x_i) $	$\text{sign}[y_i - f(x_i)]$
Regression	Huber	$y_i - f(x_i)$ for $ y_i - f(x_i) \leq \delta_m$ $\delta_m \text{sign}[y_i - f(x_i)]$ for $ y_i - f(x_i) > \delta_m$ where $\delta_m = \alpha \text{th-quantile}\{ y_i - f(x_i) \}$
Classification	Deviance	$k\text{th component: } I(y_i = \mathcal{G}_k) - p_k(x_i)$

Algorithm: GBDT

Input: $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$, where $x_i \in X$, $y_i \in \{-1, 1\}$




Initialize: $f_o(x) = \operatorname{argmin}_{\mu} \sum_{i=1}^n L(y_i, \mu)$

Output: $\hat{f}(x) = f_M(x)$

```
1  for m = 1,2, ..., M do
2      for i = 1,2, ..., n do
3          rim = - [  $\frac{\partial L(y_i, f_{m-1}(x_i))}{\partial f_{m-1}(x_i)}$  ]
4          Fit a regression tree to targets rim giving terminal regions
5          Rjm, j = 1,2, ..., Jm
6      end
7      for j = 1,2, ..., Jm do
8          μjm = argminμ  $\sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + μ)$ , j = 1,2, ..., Jm
9          Update fm(x) = fm-1(x) +  $\sum_{j=1}^{J_m} μ_j m(x \in R_{jm})$ 
10     end
11 end
```

Summary

- Instead of training different models on same data, training same model will
 - multiple times on different data sets
 - combine different models
- We can use some simple/weak model as the base model
- How do we get multiple training data sets (in practice, we only have one data set at training time)?



Thank You

SMIL内部资料 请勿外泄