

# An Introduction to Reinforcement Learning

**Prof. Mingkui Tan**

SCUT Machine Intelligence Laboratory (SMIL)



SMIL内部资料 请勿外泄

# Reinforcement Learning Applications




AlphaGo, DeepMind



Dota 2 AI, OpenAI



SMIL 内部资料请勿外泄




自动化机器学习平台, Google  
2

# AI Taught Itself to Walk

It might  
look goofy ...



# AI Learns to Park



SMIL内部资料 请勿外泄

# ChatGPT背后关键技术：人类反馈强化学习 (RLHF)

## 阶段1：监督训练

从数据集中采样问题



根据采样问题，人工给出高质量回答



基于问题和回答，对 GPT-3.5 进行监督训练



We give treats and punishments to teach...



SFT



自自自

## 阶段2：奖励模型训练（引入人类反馈）

从数据集采样问题并输入第一阶段训练的模型  
获取多条输出



人工对多个回答按照质量好坏排序



根据排序结果训练奖励模型



A In reinforcement learning, the agent is...  
B Explain rewards...  
C In machine learning...  
D We give treats and punishments to teach...



D > C > A > B



D > C > A > B

## 阶段3：基于强化学习的语言模型训练

从数据集采样新的问题



将问题输入第一阶段训练好的语言模型，根据输入生成文本



利用奖励模型给出人类价值判断，生成奖励值



使用强化学习算法，利用奖励值更新模型




Once upon a time...




r<sub>k</sub>

Prompts Dataset




Sample many prompts

Train on {sample, reward} pairs



Outputs are ranked (relative, ELO, etc.)

Initial Language Model



Generated text



Human Scoring


Prompts Dataset



x: A dog is...

$\theta$


Initial Language Model



Base Text

y: a furry mammal

Tuned Language Model (RL Policy)



Parameters Frozen\*

RLHF Tuned Text

y: man's best friend

Reinforcement Learning Update (e.g. PPO)

$$\theta \leftarrow \theta + \nabla_{\theta} J(\theta)$$



$$r_{\theta}(y|x) = -\lambda_{KL} D_{KL}(\pi_{PPO}(y|x) || \pi_{base}(y|x)) + KL prediction shift penalty$$

SMIL 内部资料 请勿外泄

# Contents

## 1 What is Reinforcement Learning?

## 2 Markov Decision Process for Reinforcement Learning

- Markov Process
- Markov Reward Process
- Markov Decision Process

## 3 Policy Gradient Methods for Reinforcement Learning

## 4 Reinforcement Learning Example: AlphaGo

## 5 Summary

# Contents

## 1 What is Reinforcement Learning?

## 2 Markov Decision Process for Reinforcement Learning

- Markov Process
- Markov Reward Process
- Markov Decision Process

## 3 Policy Gradient Methods for Reinforcement Learning

## 4 Reinforcement Learning Example: AlphaGo

## 5 Summary

SMIL内部资料 请勿外泄



# Reinforcement Learning

## ● What is Reinforcement Learning?

Learning to solve sequential decision making problems

## ● How it works?

Trial and error in a world that provides occasional rewards



SMIL内部资料 请勿外泄


# Characteristics of Reinforcement Learning

What makes reinforcement learning different from other machine learning paradigms?

- There is no supervisor, only a *reward* signal
- Feedback is delayed, not instantaneous
- Time really matters (sequential, non *i.i.d.* data)
- Agent's actions affect the subsequent data it receives

# Agent and Environment


- At each step  $t$  the agent:
  - Executes action  $A_t$
  - Receives observation  $O_t (S_t)$
  - Receives scalar reward  $R_t$
- The environment:
  - Receives action  $A_t$
  - Emits scalar reward  $R_t$
  - Emits observation  $O_{t+1}(S_{t+1})$



# Reward

- A **reward**  $R_t$  is a scalar feedback signal at step  $t$
- The agent's job is to **maximise cumulative reward**

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^k R_{t+k+1} + \dots \\ = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$



Make a humanoid robot walk



- positive  $R_t$  (+1) for moving forward
  - negative  $R_t$  (-1) for falling down
- SMI内部资料 请勿外泄

Super Mario



- positive reward  $R_t$  (+1) for getting a gold coin


# Sequential Decision Making

- Objective: Let an agent select **a series of actions** to **maximise total future rewards via some policy**:



$$\pi(a|s) = P [a_t = a|s_t = s]$$

$s$ : the current state

$a$ : possible actions given current state: e.g., move UP or DOWN



Atari: Pong



Action: move  
UP or DOWN

SMIL 内部资料 请勿外泄

# Sequential Decision Making

- The **trajectory** is the sequence of observations, actions, rewards,


$$\tau = \langle S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{t-1}, A_{t-1}, R_t \rangle$$

- State**  $S_t$  is determined *by previous trajectory* :


$$P(S_t) = P(S_t | S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{t-1}, A_{t-1}, R_t)$$

- The computation of the probability is much more complex!

**Hypothesis:** we introduce **Markov Property** to alleviate this issue!



Atari: Pong



# Basic Tasks of Reinforcement Learning

## ■ Prediction aka Evaluation

**Policy/State** Evaluation: Instant Reward

**Difficulty:** Sequential Decision Making; **Agent's actions affect the subsequent data it receives**

**(Dynamic Programming)**

$$\begin{aligned} v^\pi(s) &= \sum_{a \in A} \pi(a|s)[R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a)v^\pi(s')] \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a)[r + \gamma v^\pi(s')] \end{aligned}$$

## ■ Control: Making **Optimal** Decisions

1. 输入: MDP  $< S, A, P, R, \gamma >$

2. 输出: 最佳价值函数  $v^*(s) = \max_\pi v^\pi(s)$

以及最佳策略  $\pi^*(s) = \arg \max_\pi v^\pi(s)$

**How to make an Agent to Predict and Control?**  
SMIL内部资料 请勿外泄



RUSSIA-UKRAINE WAR

Which Ukrainian regions is Russia annexing?

Four partially Russian-controlled regions of Ukraine - Donetsk, Luhansk, Kherson and Zaporizhia - will be incorporated into Russia following 'referendums' held in the regions, the Kremlin has said.



# Contents

1 What is Reinforcement Learning?

2 Markov Decision Process for Reinforcement Learning

- Markov Process
- Markov Reward Process
- Markov Decision Process

3 Policy Gradient Methods for Reinforcement Learning


4 Reinforcement Learning Example: AlphaGo

5 Summary

# Markov Process

Given a sequence of random states  $\langle s_0, s_1, \dots, s_T \rangle$ , it satisfies **Markov Property** if and only if :

$$P(s_t | s_{t-1}) = P(s_t | s_0, \dots, s_{t-2}, s_{t-1})$$



- Once the state is known, the history can be thrown away
- The state is a sufficient statistic of the future

**Markov Process:** The future is independent of the past given the present

A **Markov Process (or Markov Chain)** is a tuple  $\langle \mathcal{S}, \mathcal{P} \rangle$

- $\mathcal{S}$  is a (finite) set of states  $\mathcal{S} = \{s_0, s_1, \dots, s_T\}$
- $\mathcal{P}$  is a **state transition** probability matrix,


$$\mathcal{P}_{ss'} = P[s_{t+1} = s' | s_t = s]$$

SMIL 内部资料 请勿外泄

# State Transition Matrix

- State transition matrix  $\mathcal{P}$  specifies  $P(s_{t+1} = s' | s_t = s)$

$$\mathcal{P} = \begin{pmatrix} P(s_1|s_1) & \cdots & P(s_N|s_1) \\ \vdots & \ddots & \vdots \\ P(s_1|s_N) & \cdots & P(s_N|s_N) \end{pmatrix}$$




$$\mathcal{S} = \{s_1, s_2, s_3, s_4\}$$

$$\mathcal{P} =$$

	$s_1$	$s_2$	$s_3$	$s_4$
$s_1$	0.1	0.2	0	0.7
$s_2$	1	0	0	0
$s_3$	0	1	0	0
$s_4$	0	0.3	0.2	0.5

# Markov Reward Process



- Markov Reward Process is a **Markov Chain + rewards**

Definition of *Markov Reward Process*


A *Markov Reward Process* is a tuple  $\langle \mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- $\mathcal{S}$  is a (finite) set of states
- $\mathcal{P}$  is a state transition probability matrix,

$$\mathcal{P}_{ss'} = P[s_{t+1} = s' | s_t = s]$$

- $\mathcal{R}$  is a reward function,  $R(s_t = s) = \mathbb{E}[r_{t+1} | s_t = s]$
- $\gamma$  is a discount factor,  $\gamma \in [0, 1]$

# Markov Decision Process



- Markov Decision Process is a Markov Reward Process + Actions

## Definition of *Markov Decision Process*

A *Markov Decision Process* is a tuple  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- $\mathcal{S}$  is a finite set of states
- $\mathcal{A}$  is a finite set of actions
- $\mathcal{P}$  is a state transition probability matrix,

$$\mathcal{P}_{ss'}^a = P[s_{t+1} = s' | s_t = s, a_t = a]$$

- $\mathcal{R}$  is a reward function,  $R(s_t = s, a_t = a) = \mathbb{E}[r_{t+1} | s_t = s, a_t = a]$
- $\gamma$  is a discount factor,  $\gamma \in [0, 1]$


# Policy

## Definition of *Policy*

A policy  $\pi$  is a distribution over actions given states:

$$\pi(a|s) = P [a_t = a | s_t = s]$$

- A policy  $\pi(a|s)$ : the agent's **behavior model**
- An MDP policy depend on the current state (not the history)
- $\pi(a|s)$  can be represented by neural networks



How to learn  $\pi(a|s)$ : **maximize total future rewards!**

# Return

Definition of *Return* (*the total future rewards*)

The return  $G_t$  is the total discounted reward from time-step  $t$

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$


- $\gamma \in [0,1]$ : discount factor for weighting the future rewards
  - $\gamma$  is used to trade off **immediate** reward and **delayed** reward
    - $\gamma$  close to 0 leads to “short-term” evaluation
    - $\gamma$  close to 1 leads to “long-term” evaluation
- Why use discount factor?
- Avoids infinite returns in cyclic Markov processes
  - Ensures the convergence when solving an MDP by dynamic programming

# State Value Function for MRP

## Definition (State-value function )

The state-value function  $V(s)$  is the expected return starting from state  $s$ :

$$V(s) = \mathbb{E}[G_t | s_t = s]$$



# Bellman Equation for MRP

- The **state-value** function can be decomposed into immediate reward  $R_{t+1}$  and discounted value of successor state  $\gamma V(s_{t+1})$

$$\begin{aligned}V(s) &= \mathbb{E}[G_t | s_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | s_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots) | s_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma G_{t+1} | s_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma V(s_{t+1}) | s_t = s]\end{aligned}$$

- Bellman equation describes the iterative relations of states

$$V(s) = R(s) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s)V(s')$$

The Bellman Equation indicates the value function of the current state can be evaluated by the next state

# Bellman Equation in Matrix Form

$$V(s) = R(s) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s)V(s')$$

- The Bellman equation can be expressed concisely using matrices

$$\begin{bmatrix} v(1) \\ \vdots \\ v(n) \end{bmatrix} = \begin{bmatrix} \mathcal{R}(1) \\ \vdots \\ \mathcal{R}(n) \end{bmatrix} + \gamma \begin{bmatrix} \mathcal{P}_{11} & \dots & \mathcal{P}_{1n} \\ \vdots & \vdots & \vdots \\ \mathcal{P}_{n1} & \dots & \mathcal{P}_{nn} \end{bmatrix} \begin{bmatrix} v(1) \\ \vdots \\ v(n) \end{bmatrix}$$

$$v = \mathcal{R} + \gamma \mathcal{P}v$$

$$(I - \gamma \mathcal{P})v = \mathcal{R}$$

$$v = (I - \gamma \mathcal{P})^{-1} \mathcal{R}$$

# Solution to MRP

## ■ Solving MRP when the model $\mathcal{P}$ is known:

$$\nu = (I - \gamma \mathcal{P})^{-1} \mathcal{R}$$

- Matrix inverse takes the **complexity  $O(N^3)$**  for  $N$  states
- Only possible for small MRPs
- The model  $\mathcal{P}$  must be known

## ■ **Iterative** methods for large MRPs:

- Dynamic Programming
- Monte-Carlo evaluation
- Temporal-Difference learning

# Contents

1 What is Reinforcement Learning?

2 Markov Decision Process for Reinforcement Learning

- Markov Process
- Markov Reward Process
- Markov Decision Process

3 Policy Gradient Methods for Reinforcement Learning

4 Reinforcement Learning Example: AlphaGo

5 Summary


# Policy of Agent

- Policy  $\pi$  can be represented by a network with parameter  $\theta$ :

$$\pi_{\theta}(a|s) = P(a|s; \theta)$$


where  $\pi_{\theta}(a_t|s)$  denotes the probability of taking an action  $a_t$  given state  $s$

- **Input:** state  $s$
- **Output:** each action  $a$  corresponds to a neuron in output layer
- Take the action based on the output probability




How to learn  $\pi_{\theta}$ ?

# Objective Function




# What is a Trajectory $\tau$ ?



# Policy Gradient

Learn the policy  $\pi_\theta(a_t|s_t)$  by **maximizing total future rewards**:

$$\max_{\theta} J(\theta) = \sum_{\tau} P(\tau; \theta)R(\tau)$$



**Policy Gradient** algorithm:

**Input:** random initialized policy  $\pi_\theta$ , max number of episodes  $N$ , learning rate  $\eta$

**Output:**  $\pi_\theta$

**for**  $i = 1$  **to**  $N$  **do**

How to compute  $\nabla_{\theta}J(\theta)$ ?

Sample a trajectory  $\tau^{(i)} = \{s_0^{(i)}, a_0^{(i)}, r_1^{(i)}, \dots, s_{T_i}^{(i)}, a_{T_i}^{(i)}, r_{T_i+1}^{(i)}\}$

**for**  $t = 0$  **to**  $T_i$  **do**

$G = \nabla_{\theta}J(\theta)$  // calculate the gradient


$\theta \leftarrow \theta + \eta G$  // maximize the objective function by ascending the gradient

**end for**

**end for**

SMILE 内部资料 请勿外泄

# Policy Gradient



SMIL内部资料 请勿外泄

# How to compute the Gradient $\nabla_{\theta}J(\theta)$ ?

The objective function is:  $\max_{\theta} J(\theta) = \sum_{\tau} P(\tau; \theta)R(\tau)$

Taking the gradient w.r.t.  $\theta$  gives

$$\nabla_{\theta}J(\theta) = \sum_{\tau} \nabla_{\theta}P(\tau; \theta)R(\tau)$$

$$\begin{aligned} &= \sum_{\tau} P(\tau; \theta) \frac{\nabla_{\theta}P(\tau; \theta)}{P(\tau; \theta)} R(\tau) \\ &= \sum_{\tau} P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta) R(\tau) \end{aligned}$$

$$\begin{aligned} \nabla f(x) &= \\ f(x) \nabla \log f(x) \end{aligned}$$

$$R(\tau) = \sum_t \gamma^t r(s_t, a_t)$$

$$P(\tau; \theta) = p(s_0) \prod_{t=0}^T \pi_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

Approximate the gradient,

$$\nabla_{\theta}J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log P(\tau^{(i)}; \theta) R(\tau^{(i)})$$

How to compute  $\nabla_{\theta} \log P(\tau^{(i)}; \theta)$ ?

# How to compute $\nabla_{\theta} \log P(\tau^{(i)}; \theta)$ ?

$$P(\tau ; \theta) = p(s_0) \prod_{t=0}^T \pi_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t)$$


$$\begin{aligned}\nabla_{\theta} \log P(\tau^{(i)}; \theta) &= \nabla_{\theta} \log \left[ \prod_{t=0}^T \underbrace{p\left(s_{t+1}^{(i)} \mid s_t^{(i)}, a_t^{(i)}\right)}_{dynamics\ model} \cdot \underbrace{\pi_{\theta}\left(a_t^{(i)} \mid s_t^{(i)}\right)}_{policy} \right] \\ &= \nabla_{\theta} \left[ \sum_{t=0}^T \log p\left(s_{t+1}^{(i)} \mid s_t^{(i)}, a_t^{(i)}\right) + \sum_{t=0}^T \log \pi_{\theta}\left(a_t^{(i)} \mid s_t^{(i)}\right) \right] \\ &= \nabla_{\theta} \sum_{t=0}^T \log \pi_{\theta}\left(a_t^{(i)} \mid s_t^{(i)}\right) \\ &= \sum_{t=0}^T \underbrace{\nabla_{\theta} \log \pi_{\theta}\left(a_t^{(i)} \mid s_t^{(i)}\right)}_{no\ dynamics\ model\ required}\end{aligned}$$

Finally, we can obtain  $\nabla_{\theta} J(\theta)$ :  $\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N R(\tau^{(i)}) \sum_{t=0}^{T_i} \nabla_{\theta} \log \pi_{\theta}\left(a_t^{(i)} \mid s_t^{(i)}\right)$

SMIL 内部资料 请勿外泄

# Policy Gradient

$$\max_{\theta} J(\theta) = \sum_{\tau} P(\tau; \theta) R(\tau)$$



**Policy Gradient** algorithm:

**Input:** random initialized policy  $\pi_\theta$ , max number of episodes  $N$ , learning rate  $\eta$

**Output:**  $\pi_\theta$

**for**  $i = 1$  **to**  $N$  **do**

    Sample a trajectory  $\tau^{(i)} = \{s_0^{(i)}, a_0^{(i)}, r_1^{(i)}, \dots, s_{T_i}^{(i)}, a_{T_i}^{(i)}, r_{T_i+1}^{(i)}\}$

**for**  $t = 0$  **to**  $T_i$  **do**

$\nabla_{\theta} J(\theta) = R(s_t^{(i)}, a_t^{(i)}, r_{t+1}^{(i)}, \dots, r_{T_i+1}^{(i)}) \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)})$

$\theta \leftarrow \theta + \eta \nabla_{\theta} J(\theta)$


**end for**

**end for**   SMIL内部资料 请勿外泄

# Differences from Gradient Descent

Consider as a classification problem

$$s_t^i \quad a_t^i \quad R(\tau^i)$$



$$\frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{T_i} \log p_\theta(a_t^i | s_t^i) \xrightarrow{\text{TF, PyTorch ...}} \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{T_i} \nabla \log p_\theta(a_t^i | s_t^i)$$

Implementation for classification:

Only difference

$$\frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{T_i} R(\tau^i) \log p_\theta(a_t^i | s_t^i) \xrightarrow{\text{TF, PyTorch ...}} \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{T_i} R(\tau^i) \nabla \log p_\theta(a_t^i | s_t^i)$$

Implementation for RL:


SMIL内部资料 请勿外泄

# Tip : Add a Baseline

$$\theta \leftarrow \theta + \eta \nabla J_\theta$$

It is possible that  $R(\tau^i)$  is always positive

$$\nabla_\theta J \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{T_i} (R(\tau^i) - b) \nabla \log p_\theta(a_t^i | s_t^i) \quad b \approx E[R(\tau)]$$



SMIL内部资料 请勿外泄

# Contents

1 What is Reinforcement Learning?

2 Markov Decision Process for Reinforcement Learning

- Markov Process
- Markov Reward Process
- Markov Decision Process


3 Policy Gradient Methods for Reinforcement Learning

4 Reinforcement Learning Example: AlphaGo

5 Summary

# What is Go?

- An abstract board game for two players
- Played on 19 x 19 board
- Playing pieces are black and white stones
- Stones placed on vacant intersections of board
- Player which surrounds more territory wins



SMIL内部资料 请勿外泄

# Challenges for AI in Cracking Go

- Impossible to calculate every possible move on board
- Brute-force method used by most AIs clearly fails
  - Search space is **huge**
  - Impossible for computers to evaluate who is winning
- Go requires **more intuition and experience** than just logic
- Becomes necessary to mimic human mind


# AlphaGo

- A computer program that plays Go game
- Developed by Google DeepMind in 2016
- Not a pre-programmed algorithm
- Can actually learn from itself
- Reduce search space
  - Reducing action candidates
  - Board Evaluation


# Reducing Action Candidates

- Imitating expert's moves (supervised learning)

Current Board



Next Action



**Expert Moves Imitator Model  
(w/ CNN)**

**Training:**  $\Delta\sigma \propto \frac{\partial \log p_\sigma(a|s)}{\partial \sigma}$

# Reducing Action Candidates

- Improving through self-plays (reinforcement learning)



SMIL内部资料 请勿外泄



# Reward

- The reward function  $r(s)$  that is zero for all non-terminal state
- The reward signal is delayed
- The outcome  $z_t = \pm r(s_T)$  is the terminal reward at the end of the game from the perspective of the current player at time step  $t$ 
  - +1 for winning
  - -1 for losing
- Rollout is inefficient

# Board Evaluation


- Estimate a value function that predicts the outcome from the position
  - Train the value network by regression on state-outcome pairs  $(s, z)$
  - Minimize the mean squared error (MSE) between the predicted value  $v_\theta(s)$ , and the corresponding outcome  $z$

$$\Delta\theta \propto \frac{\partial v_\theta(s)}{\partial \theta} (z - v_\theta(s))$$



SMIL 内部资料 请勿外泄

# Monte Carlo Tree Search in AlphaGo




- Selecting the edge with maximum action value  $Q$
- The leaf node may be expanded
- At the end of a simulation, the leaf node is evaluated by value network
- Action value  $Q$  are updated to track the mean value of all evaluations  $r(\cdot)$  and  $v_\theta(\cdot)$  in the subtree below that action

SMIL内部资料 请勿外泄


# Learning Pipeline of AlphaGo

- A fast rollout policy  $p_\pi$  and supervised learning (SL) policy network  $p_\sigma$  are trained to predict human expert moves in a data set of positions.
- A reinforcement learning (RL) policy  $p_\rho$  is initialized to the SL policy
- Then  $p_\sigma$  is improved by policy gradient learning to **maximize the outcome** (i.e., winning more game) **against previous versions of the policy network**.
- A new dataset is generated by playing games of **self-play** with the RL policy network.
- A value network  $v_\theta$  is trained by regression to predict the expected outcome




SMIL内部资料 请勿外泄

# Playing Go Using Learnt Policy



SMIL内部资料 请勿外泄

# Performance



- Results of a tournament between different Go programs. Each program used approximately 5s computation time per move.
- Performance of AlphaGo for different combinations of components.

SMIL内部资料 请勿外泄

# Contents

## 1 What is Reinforcement Learning?

## 2 Markov Decision Process for Reinforcement Learning

- Markov Process
- Markov Reward Process
- Markov Decision Process


## 3 Policy Gradient Methods for Reinforcement Learning

## 4 Reinforcement Learning Example: AlphaGo

## 5 Summary

SMIL内部资料 请勿外泄

# Branches of Machine Learning



SMIL内部资料 请勿外泄


# Categorizing of RL Agents (1)

- Value Based
  - No Policy (Implicit)
  - Value Function
- Policy Based
  - Policy
  - No Value Function
- Actor Critic
  - Policy
  - Value Function

# Categorizing of RL Agents (2)

- Model Free
  - Policy and/or Value Function
  - No Model
- Model Based
  - Policy and/or Value Function
  - Model

# Taxonomy of RL Algorithms




SMIL内部资料 请勿外泄

# Types of RL algorithms

Better  
Sample Efficient

Less  
Sample Efficient



## Model-based

- Learn the model of the world, then plan using the model
- Update model often
- Re-plan often

## Value-based

- Learn the state or state-action value
- Act by choosing best action in state
- Exploration is a necessary add-on

## Policy-based

- Learn the stochastic policy function that maps state to action
- Act by sampling policy
- Exploration is baked in

# Reference

- [1] Sutton, Richard S., et al. "Policy gradient methods for reinforcement learning with function approximation." NeurIPS. 2000.
- [2] Mnih, Volodymyr, et al. "Asynchronous methods for deep reinforcement learning." ICML. 2016.
- [3] Schulman, John, et al. "Proximal policy optimization algorithms." arXiv. 2017.
- [4] Schulman, John, et al. "Trust region policy optimization." ICML. 2015.
- [5] Lillicrap, Timothy P., et al. "Continuous control with deep reinforcement learning." ICLR.2016.
- [6] Fujimoto, Scott, Herke Van Hoof, and David Meger. "Addressing function approximation error in actor-critic methods." ICLR. 2018.
- [7] Haarnoja, Tuomas, et al. "Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor." ICML. 2018.
- [8] Mnih, Volodymyr, et al. "Playing atari with deep reinforcement learning." arXiv. 2013.

# Reference

- [9] Bellemare, Marc G., et al. "A Distributional Perspective on Reinforcement Learning." ICML. 2017.
- [10] Dabney, Will, et al. "Distributional Reinforcement Learning With Quantile Regression." AAAI. 2018.
- [11] Andrychowicz, Marcin, et al. "Hindsight experience replay." NeurIPS. 2017.
- [12] Ha, David, et al. "World models." NeurIPS. 2018.
- [13] Racanière, Sébastien, et al. "Imagination-augmented agents for deep reinforcement learning." NeurIPS. 2017.
- [14] Bansal, Somil, et al. "Mbmf: Model-based priors for model-free reinforcement learning." arXiv. 2017.
- [15] Feinberg, Vladimir, et al. "Model-Based Value Estimation for Efficient Model-Free Reinforcement Learning." arXiv. 2018.
- [16] Silver, David, et al. "Mastering the game of go without human knowledge." nature. 2017.

# Reference

- Silver, D.(2015). UCL Course on RL.
- Levine, S.(2017). CS 294: Deep Reinforcement Learning.
- Sutton, Richard S., and Andrew G. Barto. Reinforcement learning: An introduction. MIT press, 1998.
- Hung-yi Le. Deep Reinforcement Learning: Scratching the surface, 2016

# Q&A

SMIL内部资料 请勿外泄