

Overfitting, Regularization and Support Vector Machine

Prof. Mingkui Tan

SCUT Machine Intelligence Laboratory (SMIL)



SMIL内部资料 请勿外泄

Contents

1 Regression Revisited

2 Gradient Decent

Contents

1 Regression Revisited

2 Gradient Decent

SMIL内部资料 请勿外泄

Basic Concepts about Machine Learning

- Machine Learning compose of three parts:
 - Data
 - Model(function)
 - Loss(prediction)

Task of Machine Learning

Machine Learning \approx Looking for a Function

- Speech Recognition

$$f\left(\begin{array}{c} \text{[Speaker icon]} \\ \text{[Waveform plot]} \end{array} \right) = \text{“How are you”}$$

- Image Recognition

$$f\left(\begin{array}{c} \text{[Speaker icon]} \\ \text{[Image of a kitten]} \end{array} \right) = \text{“Cat”}$$

- Playing Go

$$f\left(\begin{array}{c} \text{[Speaker icon]} \\ \text{[Image of a Go board]} \end{array} \right) = \text{“5-5” (next move)}$$

- Dialogue System

$$f\left(\begin{array}{c} \text{“Hi”} \\ \text{(what the user said)} \end{array} \right) = \text{“Hello”} \quad \begin{array}{c} \text{(system response)} \end{array}$$

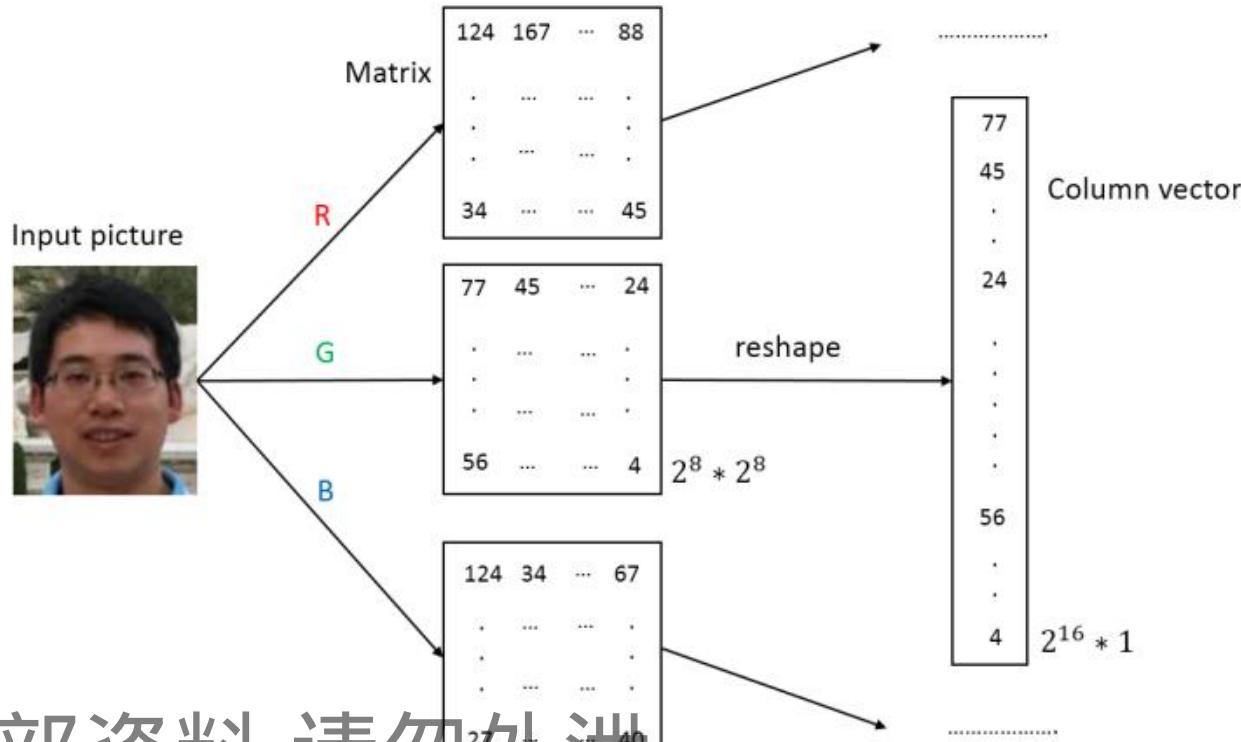
SMIL 内部资料 请勿外泄

Column Vector

■ Data:

$$D = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^n$$

- \mathbf{x} is input, and we usually present it as column vector
- For example, \mathbf{x} may be a picture stored as a matrix:
- Loss(prediction)



SMIL内部资料 请勿外泄

Regression

- We want to learn a function $f(x)$ to predict x by

$$\hat{y} = f(x)$$

- The prediction may be inconsistent with the ground truth.
- We measure the differences by some losses $l(f(x), y_i) \geq 0$:

- Absolute value loss:

$$l(f(x), y_i) = |f(x) - y_i|$$

- Least squares loss:

$$l(f(x), y_i) = \frac{1}{2}(f(x) - y_i)^2$$

Regression

- The loss function \mathcal{L}_D plays a major role in machine learning
- The smaller value of \mathcal{L}_D , the better

Find the best $\hat{\theta}$ by solving the following optimization problem:

$$f^* = \min_f \sum_{n=1}^n l(f(x), y_i)$$

Linear Regression

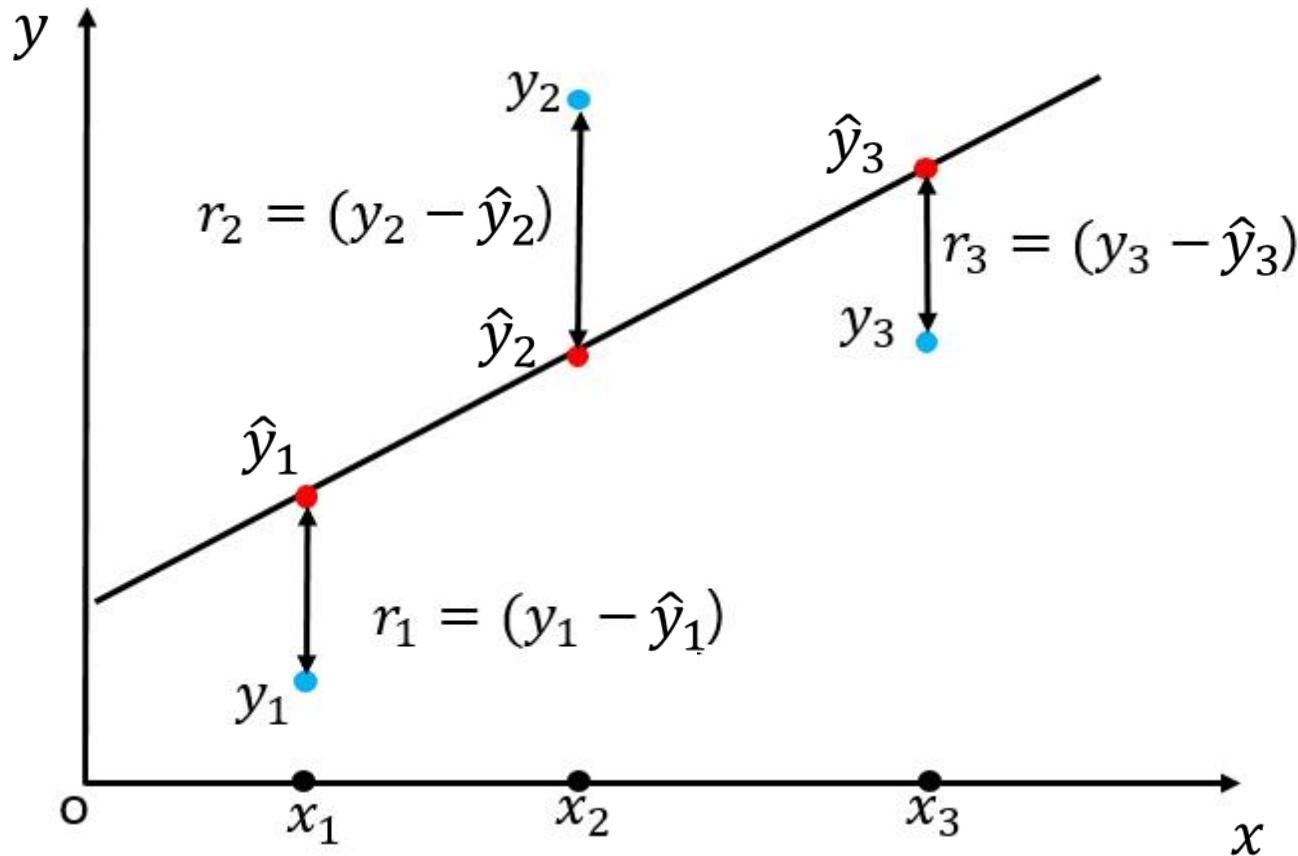
Learn $f(\mathbf{x}; \mathbf{w}, b)$ with

- Parameters: $\mathbf{w} \in \mathbb{R}^m, b \in \mathbb{R}$
- Input: \mathbf{x} where $x_i \in \mathbb{R}$, features for $i \in \{1, \dots, m\}$
- Model Function:

$$\begin{aligned}f(\mathbf{x}; \mathbf{w}, b) &= w_1 x_1 + \dots + w_m x_m + b \\&= \sum_{i=1}^m w_i x_i + b \\&= \mathbf{w}^T \mathbf{x} + b\end{aligned}$$

Linear Regression

■ What makes a good model?



Least Square Regression

■ Least squared loss

$$\begin{aligned}\mathcal{L}_D(\mathbf{w}) &= + \frac{1}{2} \sum_{i=1}^n (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 \\ &= \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2\end{aligned}$$

Training: find minimizer of least squared loss

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \mathcal{L}_D(\mathbf{w})$$

Closed-form Solution

- First-order condition of the optimal solution:

$$\frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}} = 0$$

- For the Least Regression problem, we have

$$\begin{aligned}\frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}} &= -\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \mathbf{w} = 0 \\ &\Rightarrow \mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y} \\ &\Rightarrow \mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}\end{aligned}$$

- We obtain the optimal \mathbf{w}^* by

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Issue of the Closed-form Solution

- Closed-form solution: $\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
 - The matrix $(\mathbf{X}^T \mathbf{X})^{-1}$ may not be invertible, which means the matrix may have infinite number of solutions!

Regularized Least Square (RLS) Regression

- Impose regularization on \mathbf{w} :

$$\begin{aligned}\mathcal{L}_D(\mathbf{w}) &= \frac{\lambda}{2} \|\mathbf{w}\|_2^2 + \frac{1}{2} \sum_{i=1}^n (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 \\ &= \frac{\lambda}{2} \|\mathbf{w}\|_2^2 + \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2\end{aligned}$$

- Here, $\frac{\lambda}{2} \|\mathbf{w}\|_2^2$ is called **Regularizer**, λ is called **trade-off parameter** or **regularization parameter**

Training: find minimizer of least squared loss

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \mathcal{L}_D(\mathbf{w})$$

Closed-form Solution for Regularized Least Square(RLS)

- First-order condition of the optimal solution:

$$\frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}} = 0$$

- For the Least Regression problem, we have

$$\begin{aligned}\frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}} &= \lambda \mathbf{w} - \mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \mathbf{w} = 0 \\ &\Rightarrow (\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X}) \mathbf{w} = \mathbf{X}^T \mathbf{y} \\ &\Rightarrow \mathbf{w} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}\end{aligned}$$

- We obtain the optimal \mathbf{w}^* by

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

Issue of the Closed-form Solution

- Closed-form solution: $\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$
 - The inverse of a large matrix needs huge memory
 - The inverse takes $O(m^3)$ complexity to compute

Contents

1 Regression Revisited

2 Gradient Decent

Machine Learning

- Training Procedure:
 - Identify a set of hypotheses $f(x; w)$
 - Define a loss criterion \mathcal{L}_D
 - Pick the best w^* by minimizing a loss function $\mathcal{L}_D(w)$
- Learning is done through optimization

Main Tool: Gradient

- Typical case (with possibly parameterized g):

$$\mathcal{L}_D(\mathbf{w}) : \mathbb{R}^n \mapsto \mathbb{R}$$

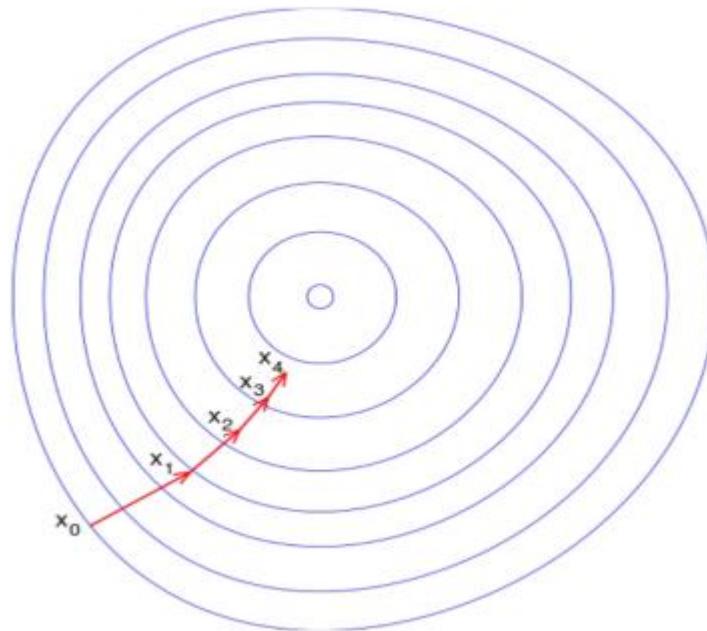
- Gradient (vector of partial derivatives)

$$\frac{\partial \mathcal{L}_D(\mathbf{w})}{\partial \mathbf{w}} = \begin{bmatrix} \frac{\partial \mathcal{L}_D(w_1)}{\partial w_1} \\ \frac{\partial \mathcal{L}_D(w_2)}{\partial w_2} \\ \vdots \\ \frac{\partial \mathcal{L}_D(w_n)}{\partial w_n} \end{bmatrix}$$

(We will always write as column vectors)

Decent Direction

- We use $\mathbf{d} = -\frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}}$ as the direction of optimization



- Why $\mathcal{L}_{\mathcal{D}}(\mathbf{w}') = \mathcal{L}_{\mathcal{D}}(\mathbf{w} + \eta \mathbf{d}) \leq \mathcal{L}_{\mathcal{D}}(\mathbf{w}), \eta \rightarrow 0^+$?

Descent Direction

Proof:

By Taylor expansion, when $\eta \rightarrow 0^+$:

$$\begin{aligned}\mathcal{L}_{\mathcal{D}}(\mathbf{w} + \eta \mathbf{d}) &= \mathcal{L}_{\mathcal{D}}(\mathbf{w}) + \left(\frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}} \right)^T \eta \mathbf{d} + o(\eta \mathbf{d}) \\ &= \mathcal{L}_{\mathcal{D}}(\mathbf{w}) + \eta' \left(\frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}} \right)^T \mathbf{d}\end{aligned}$$

Note that $\eta' > 0$ and

$$\eta' \left(\frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}} \right)^T \mathbf{d} = -\eta' \mathbf{d}^T \mathbf{d} \leq 0$$

We have:

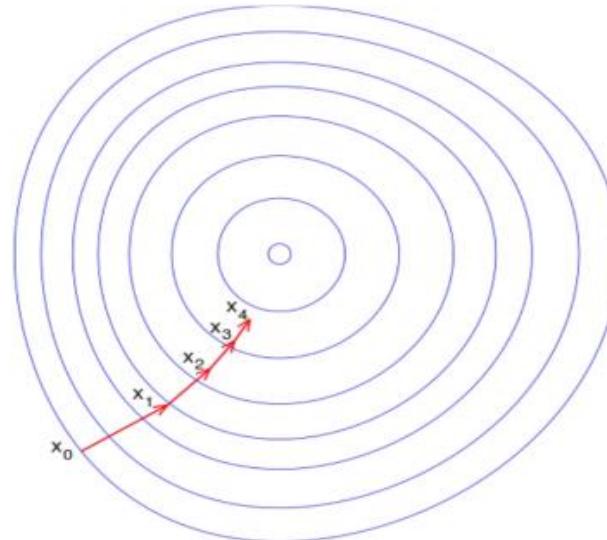
$$\mathcal{L}_{\mathcal{D}}(\mathbf{w}') = \mathcal{L}_{\mathcal{D}}(\mathbf{w} + \eta \mathbf{d}) \leq \mathcal{L}_{\mathcal{D}}(\mathbf{w})$$

Gradient Descent

Minimize loss by repeated gradient steps (when no closed form):

- Compute gradient of loss with respect to parameters $\frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}}$
- Update parameters with learning rate η

$$\mathbf{w}' = \mathbf{w} - \eta \frac{\partial \mathcal{L}_{\mathcal{D}}(\mathbf{w})}{\partial \mathbf{w}}$$



SMIL 内部资料 请勿外泄
Figure: Gradient steps on a simple $m=2$ loss function.

Appropriate Value of Learning Rate

Learning rate η has a large impact on convergence

- Too large $\eta \Rightarrow$ oscillate and may even diverge
- Too small $\eta \Rightarrow$ too slow to converge

Adaptive learning rate (For example) :

- Set larger learning rate at the beginning
- Use relatively smaller learning rate in the later epochs
- Decrease the learning rate:

$$\eta_{k+1} = \frac{\eta_k}{k + 1}$$

Thank You

SMIL内部资料 请勿外泄